

```
In [1]: # Classification Project: Rain in Australia
# The dataset from Kaggle contains daily weather observations from various locations in Australia
# Features include temperature, rainfall, sunshine, wind gusts, humidity
# The goal of this project is to build three models (Classification Trees,
```

```
In [2]: import pandas as pd
import os
import numpy as np
from tqdm import tqdm
import random
import csv
import statistics
import statsmodels.api as sm

import matplotlib
import matplotlib.pyplot as plt
import matplotlib.lines as mlines
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
from sklearn import metrics
from sklearn.preprocessing import scale
from sklearn.metrics import mean_squared_error
from sklearn import tree
from sklearn.model_selection import cross_validate
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.metrics import recall_score, f1_score
from sklearn.metrics import roc_curve, auc
```

```
In [3]: # Data loading and initial inspection, understanding the structure of data

rainAUS = pd.read_csv("/Users/aleksandra/Desktop/MachineLearning2/weather.csv")
print(rainAUS.head())
print(rainAUS.dtypes.value_counts())
```

```

          Date Location MinTemp MaxTemp Rainfall Evaporation Sunshine
\ 0 2008-12-01 Albury 13.4 22.9 0.6 NaN NaN
  1 2008-12-02 Albury 7.4 25.1 0.0 NaN NaN
  2 2008-12-03 Albury 12.9 25.7 0.0 NaN NaN
  3 2008-12-04 Albury 9.2 28.0 0.0 NaN NaN
  4 2008-12-05 Albury 17.5 32.3 1.0 NaN NaN

      WindGustDir WindGustSpeed WindDir9am ... Humidity9am Humidity3pm \
0 W 44.0 W ... 71.0 22.0
1 WNW 44.0 NNW ... 44.0 25.0
2 WSW 46.0 W ... 38.0 30.0
3 NE 24.0 SE ... 45.0 16.0
4 W 41.0 ENE ... 82.0 33.0

Pressure9am Pressure3pm Cloud9am Cloud3pm Temp9am Temp3pm RainTod
ay \
0 1007.7 1007.1 8.0 NaN 16.9 21.8
No
1 1010.6 1007.8 NaN NaN 17.2 24.3
No
2 1007.6 1008.7 NaN 2.0 21.0 23.2
No
3 1017.6 1012.8 NaN NaN 18.1 26.5
No
4 1010.8 1006.0 7.0 8.0 17.8 29.7
No

RainTomorrow
0 No
1 No
2 No
3 No
4 No

[5 rows x 23 columns]
float64 16
object 7
Name: count, dtype: int64

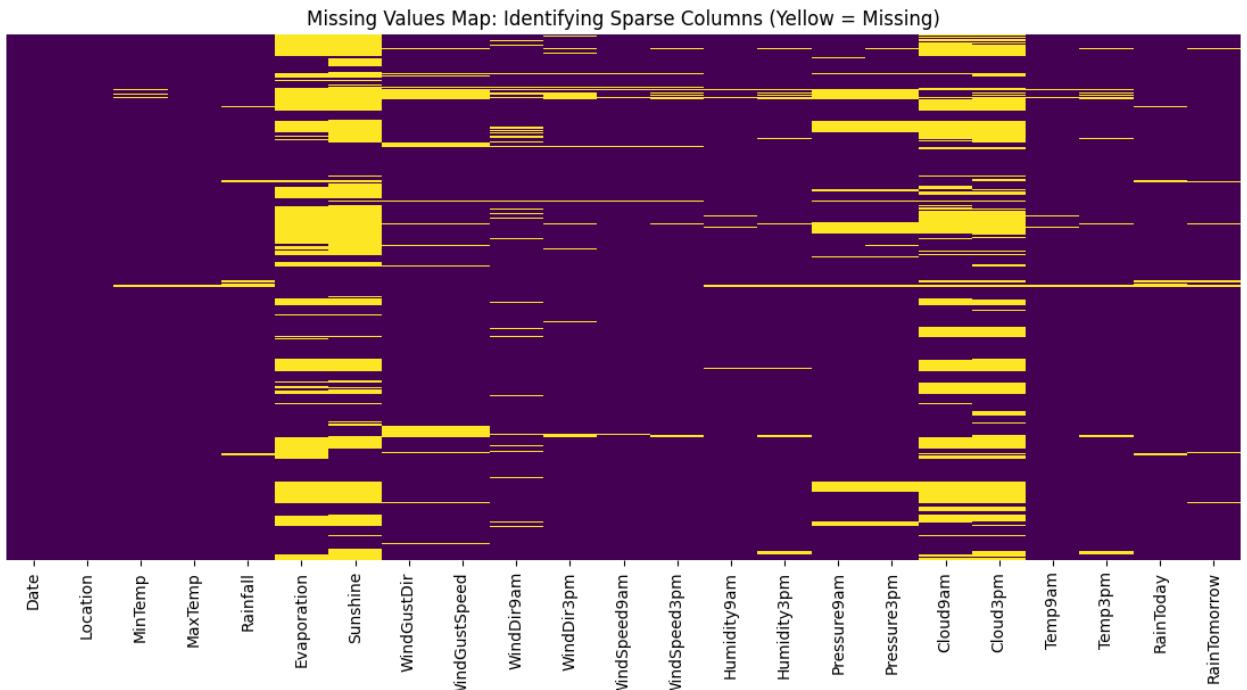
```

In [4]: #Variables with very high proportions of missing values (above 40%) are e

```

plt.figure(figsize=(14, 6))
# Heatmap of missing values: Yellow lines indicate missing data
sns.heatmap(rainAUS.isnull(), cbar=False, yticklabels=False, cmap='viridis')
plt.title('Missing Values Map: Identifying Sparse Columns (Yellow = Missing)')
plt.show()
rainAUS.isna().mean() * 100
rainAUS = rainAUS.drop(columns=["Evaporation", "Sunshine", "Cloud9am", "C"])
print(rainAUS.head())

```



	Date	Location	MinTemp	MaxTemp	Rainfall	WindGustDir	WindGustSp
eed \ 4.0	2008-12-01	Albury	13.4	22.9	0.6	W	4
4.0	2008-12-02	Albury	7.4	25.1	0.0	NNW	4
6.0	2008-12-03	Albury	12.9	25.7	0.0	WSW	4
4.0	2008-12-04	Albury	9.2	28.0	0.0	NE	2
1.0	2008-12-05	Albury	17.5	32.3	1.0	W	4

	WindDir9am	WindDir3pm	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity
3pm \ 2.0	W	NNW	20.0	24.0	71.0	2
5.0	NNW	WSW	4.0	22.0	44.0	2
0.0	W	WSW	19.0	26.0	38.0	3
6.0	SE	E	11.0	9.0	45.0	1
3.0	ENE	NW	7.0	20.0	82.0	3

	Pressure9am	Pressure3pm	Temp9am	Temp3pm	RainToday	RainTomorrow
0	1007.7	1007.1	16.9	21.8	No	No
1	1010.6	1007.8	17.2	24.3	No	No
2	1007.6	1008.7	21.0	23.2	No	No
3	1017.6	1012.8	18.1	26.5	No	No
4	1010.8	1006.0	17.8	29.7	No	No

In [5]: # Data type distribution and missing values

```
print("Frequency of data type")
rainAUS.dtypes.value_counts()
print(rainAUS.head())
print(rainAUS.isna().sum())
rainAUS = rainAUS.dropna(subset=["RainTomorrow"])
```

Frequency of data type

	Date	Location	MinTemp	MaxTemp	Rainfall	WindGustDir	WindGustSp
eed \ 4.0	2008-12-01	Albury	13.4	22.9	0.6	W	4
4.0	2008-12-02	Albury	7.4	25.1	0.0	NNW	4
6.0	2008-12-03	Albury	12.9	25.7	0.0	WSW	4
4.0	2008-12-04	Albury	9.2	28.0	0.0	NE	2
1.0	2008-12-05	Albury	17.5	32.3	1.0	W	4

	WindDir9am	WindDir3pm	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity
3pm \ 2.0	W	NNW	20.0	24.0	71.0	2
5.0	NNW	WSW	4.0	22.0	44.0	2
0.0	W	WSW	19.0	26.0	38.0	3
6.0	SE	E	11.0	9.0	45.0	1
3.0	ENE	NW	7.0	20.0	82.0	3

	Pressure9am	Pressure3pm	Temp9am	Temp3pm	RainToday	RainTomorrow
0	1007.7	1007.1	16.9	21.8	No	No
1	1010.6	1007.8	17.2	24.3	No	No
2	1007.6	1008.7	21.0	23.2	No	No
3	1017.6	1012.8	18.1	26.5	No	No
4	1010.8	1006.0	17.8	29.7	No	No

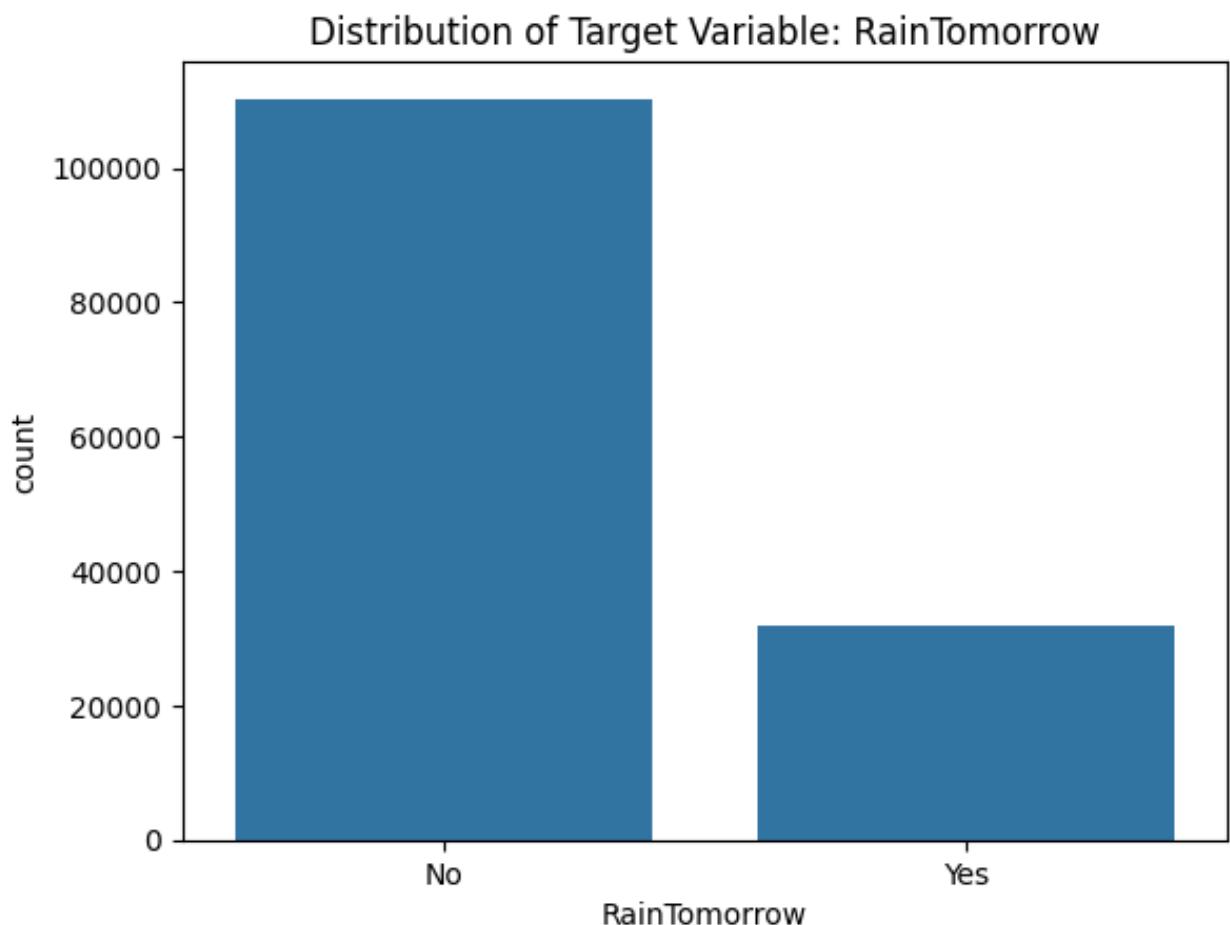
Date 0
 Location 0
 MinTemp 1485
 MaxTemp 1261
 Rainfall 3261
 WindGustDir 10326
 WindGustSpeed 10263
 WindDir9am 10566
 WindDir3pm 4228
 WindSpeed9am 1767
 WindSpeed3pm 3062
 Humidity9am 2654
 Humidity3pm 4507
 Pressure9am 15065
 Pressure3pm 15028
 Temp9am 1767
 Temp3pm 3609
 RainToday 3261
 RainTomorrow 3267
 dtype: int64

```
In [6]: # Target variable analysis - distribution
import seaborn as sns
import matplotlib.pyplot as plt

sns.countplot(x="RainTomorrow", data=rainAUS)
plt.title("Distribution of Target Variable: RainTomorrow")
plt.show()

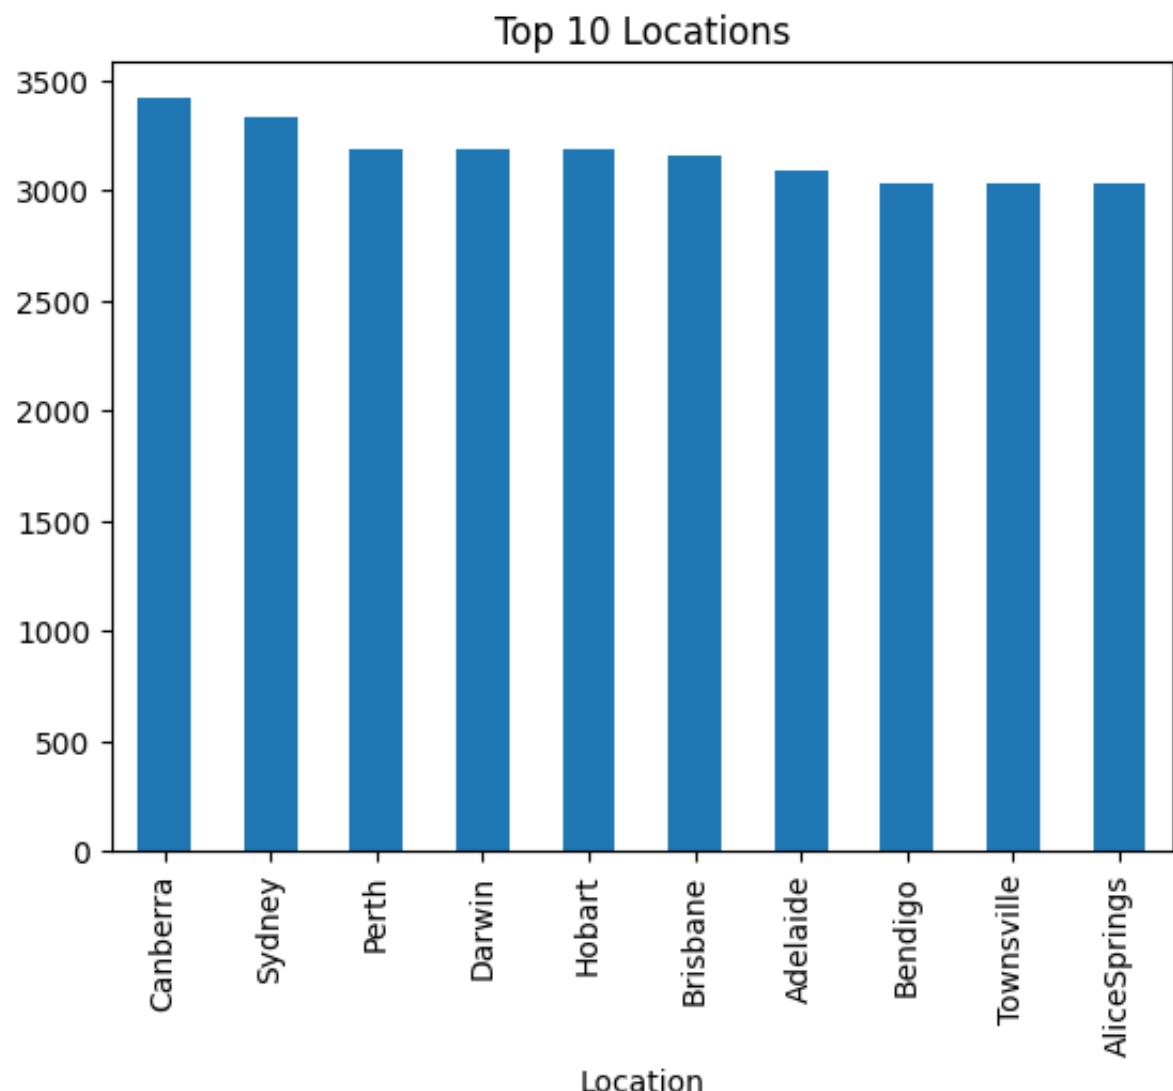
# Percentage distribution

rainAUS[["RainTomorrow"]].value_counts(normalize=True) * 100
```

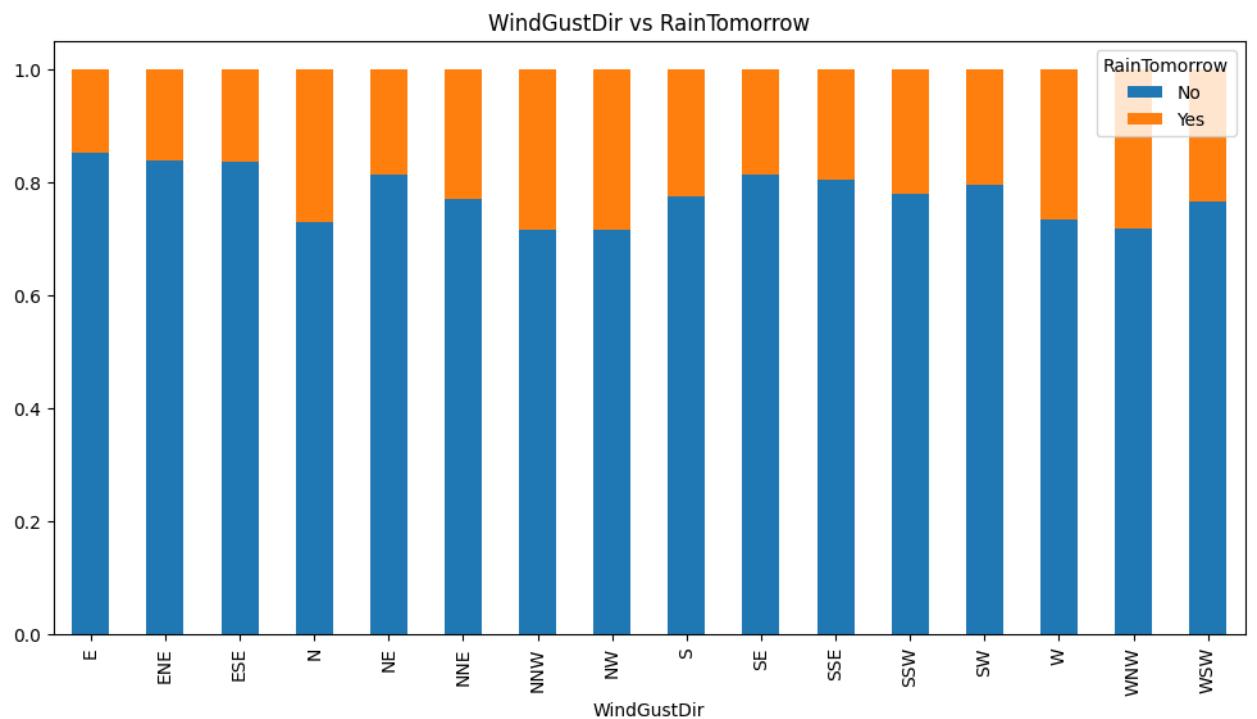


```
Out[6]: RainTomorrow
No      77.581878
Yes     22.418122
Name: proportion, dtype: float64
```

```
In [7]: rainAUS[["Location"]].value_counts().head(10).plot(kind="bar")
plt.title("Top 10 Locations")
plt.show()
```



```
In [8]: pd.crosstab(rainAUS["WindGustDir"], rainAUS["RainTomorrow"], normalize="index", kind="bar", stacked=True, figsize=(12,6))
plt.title("WindGustDir vs RainTomorrow")
plt.show()
```

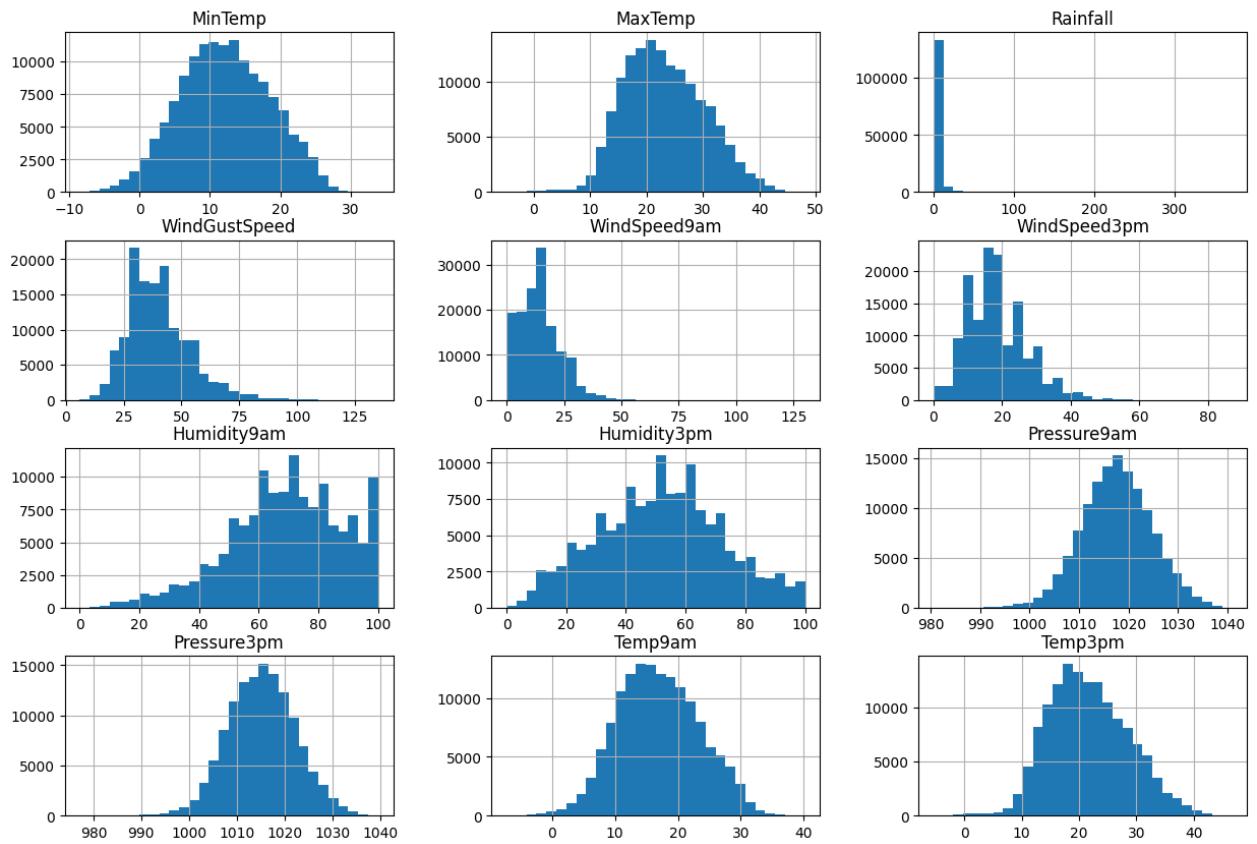


```
In [9]: num_cols = rainAUS.select_dtypes(include="number").columns

rainAUS[num_cols].hist(bins=30, figsize=(15,10))
plt.suptitle("Distributions of Numerical Features")
plt.show()

# Temperature and pressure variables exhibit approximately normal distribution
# strongly right-skewed, reflecting the presence of rare extreme events.
# indicating potential predictive relevance for rainfall occurrence. The
# skewness motivate the use of tree-based classification models, which are
```

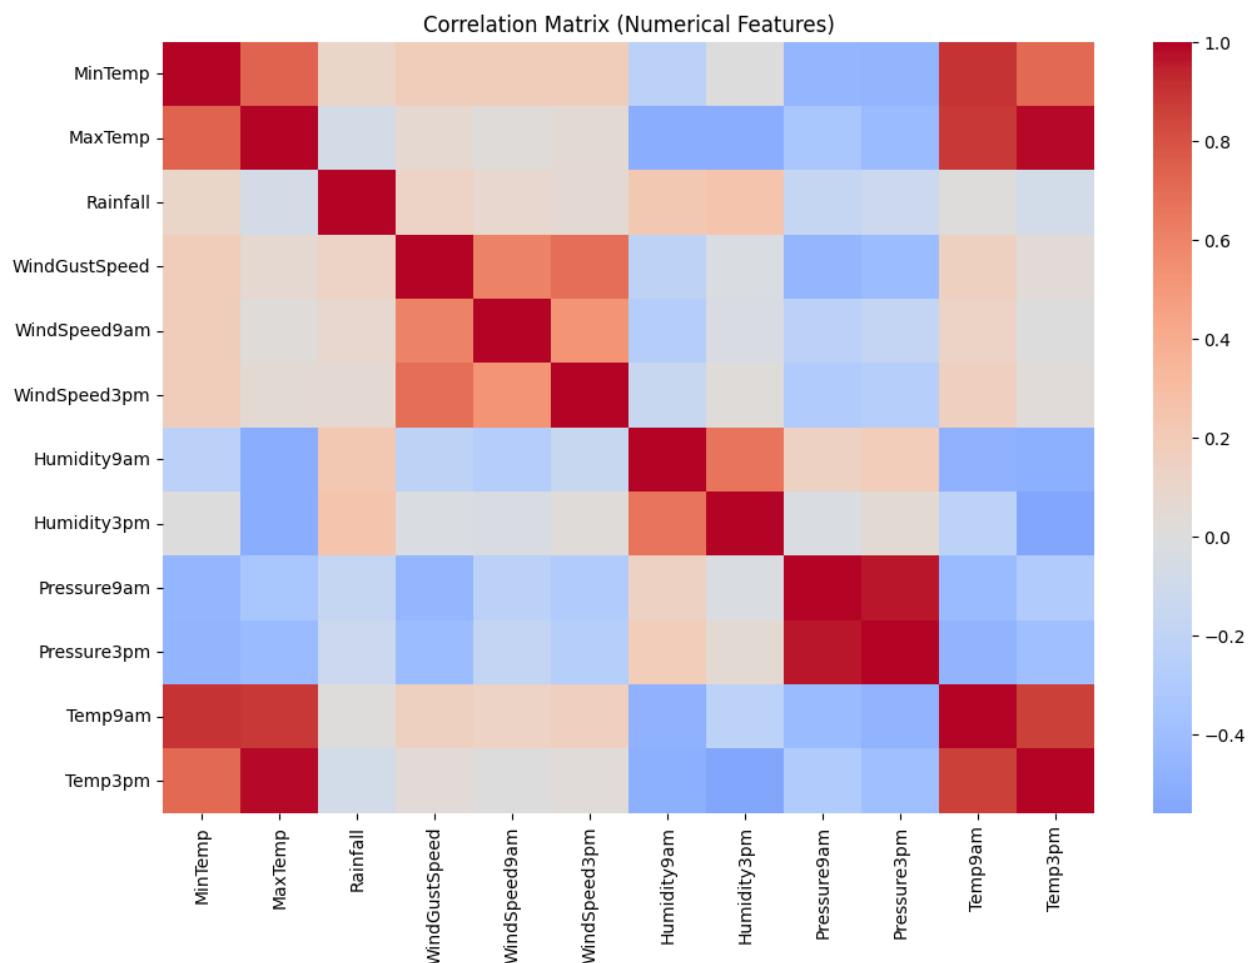
Distributions of Numerical Features



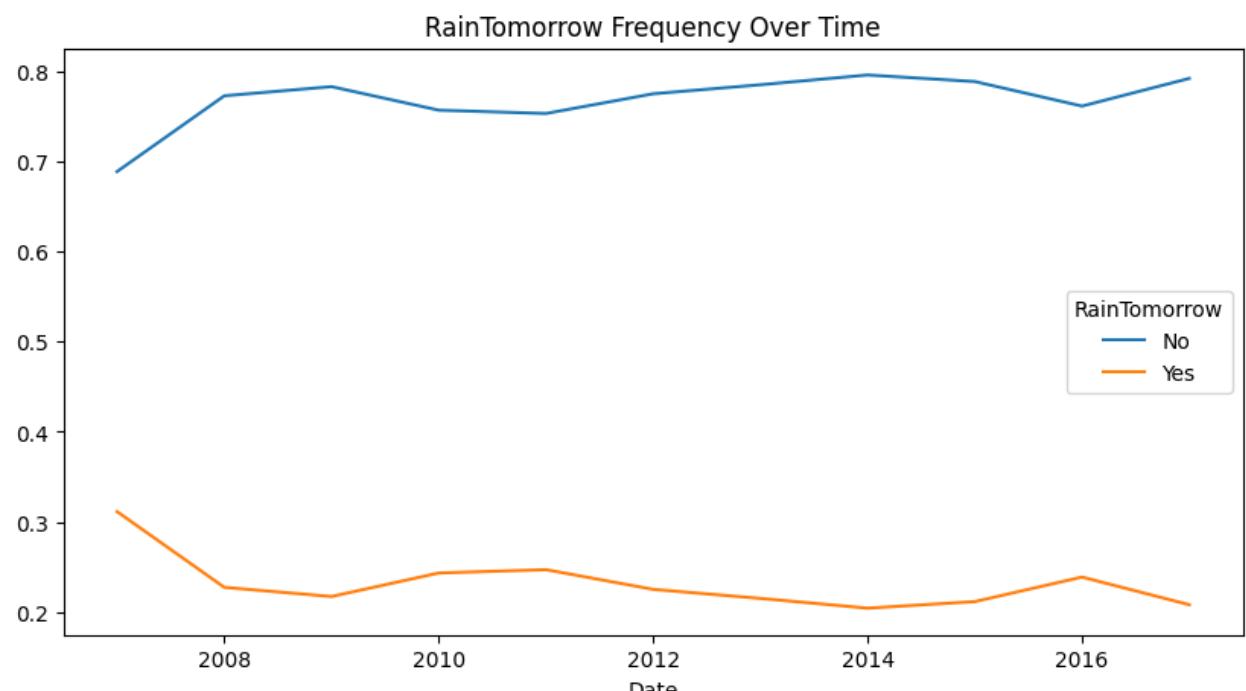
```
In [10]: num_cols = rainAUS.select_dtypes(include="number").columns

plt.figure(figsize=(12,8))
sns.heatmap(
    rainAUS[num_cols].corr(),
    cmap="coolwarm",
    center=0
)
plt.title("Correlation Matrix (Numerical Features)")
plt.show()

# Temperature and Humidity - negative correlations (expected meteorologic
```

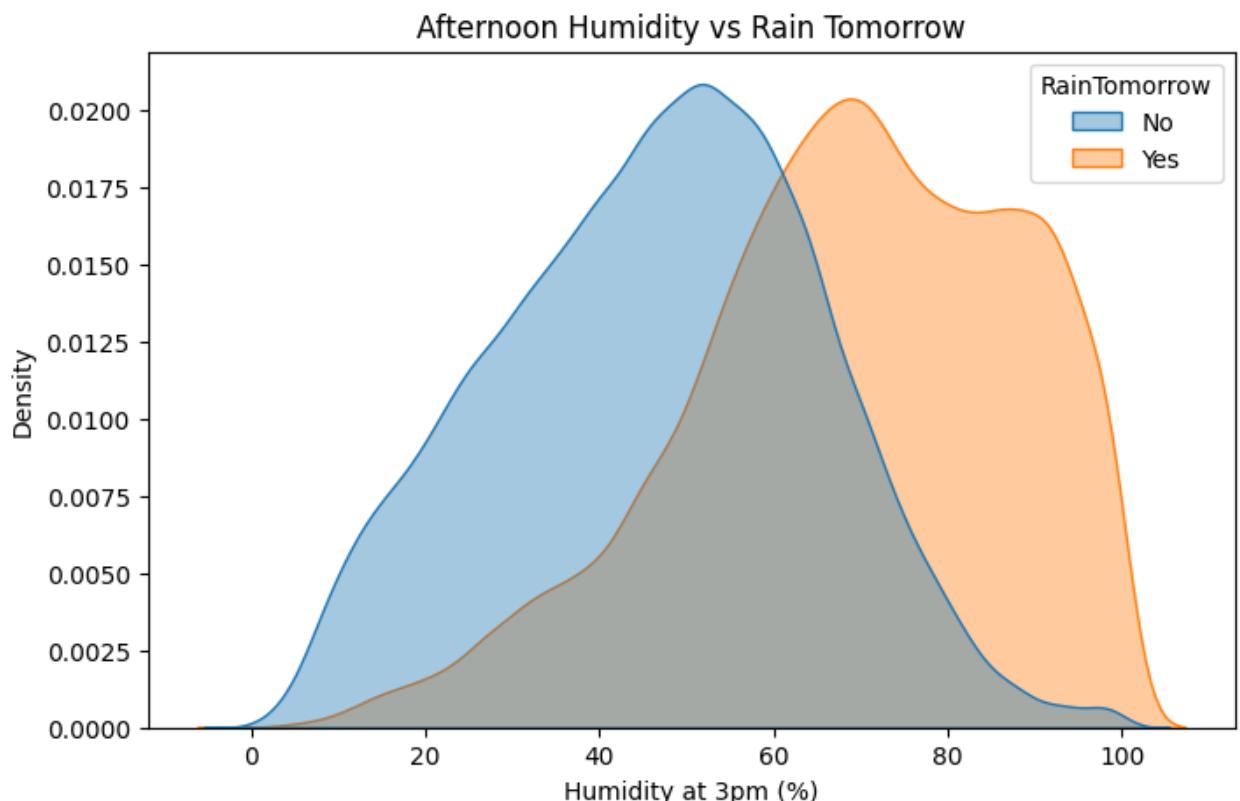


```
In [11]: rainAUS["Date"] = pd.to_datetime(rainAUS["Date"])
rainAUS.groupby(rainAUS["Date"].dt.year)[["RainTomorrow"]].value_counts(normalize=True).unstack().plot(figsize=(10,5))
plt.title("RainTomorrow Frequency Over Time")
plt.show()
```



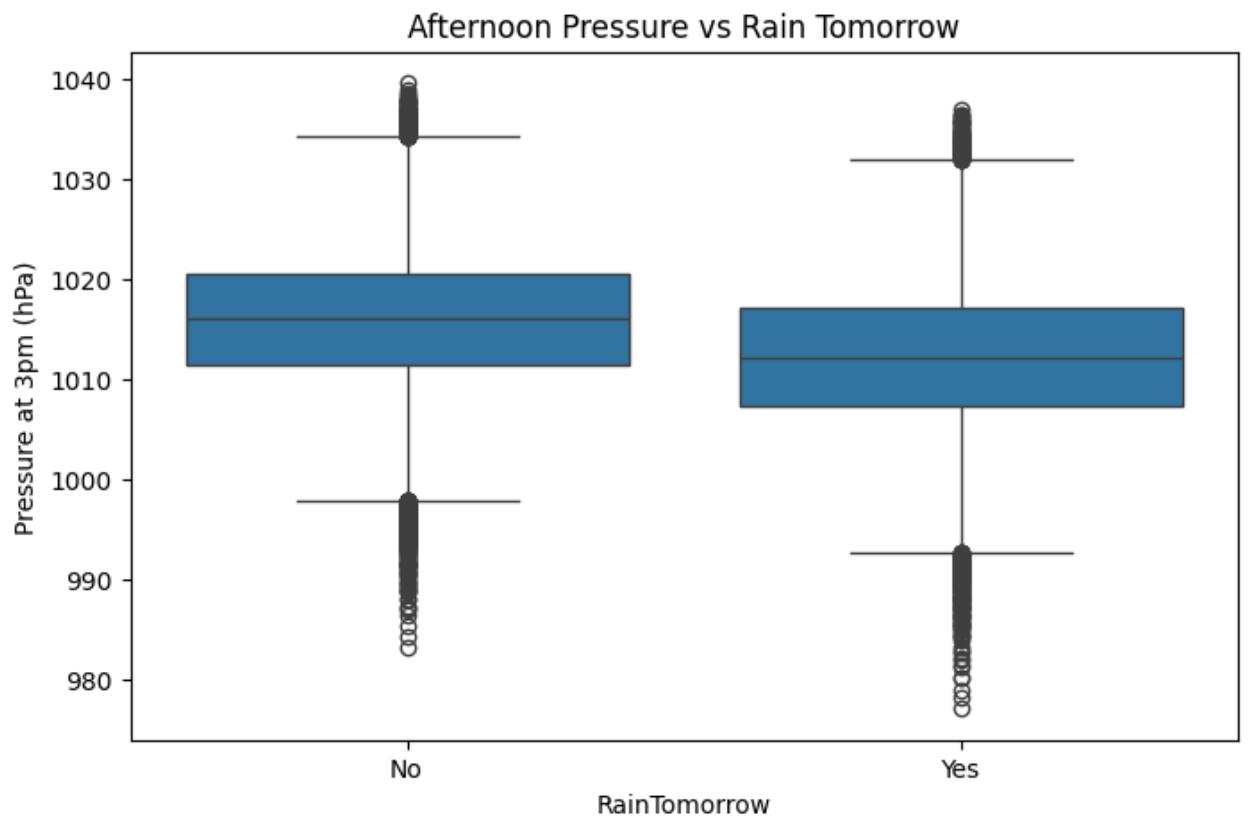
```
In [12]: plt.figure(figsize=(8,5))
sns.kdeplot(
    data=rainAUS,
    x="Humidity3pm",
    hue="RainTomorrow",
    common_norm=False,
    fill=True,
    alpha=0.4
)
plt.title("Afternoon Humidity vs Rain Tomorrow")
plt.xlabel("Humidity at 3pm (%)")
plt.show()

#Rain rarely comes suddenly – the atmosphere is already saturated the day
```



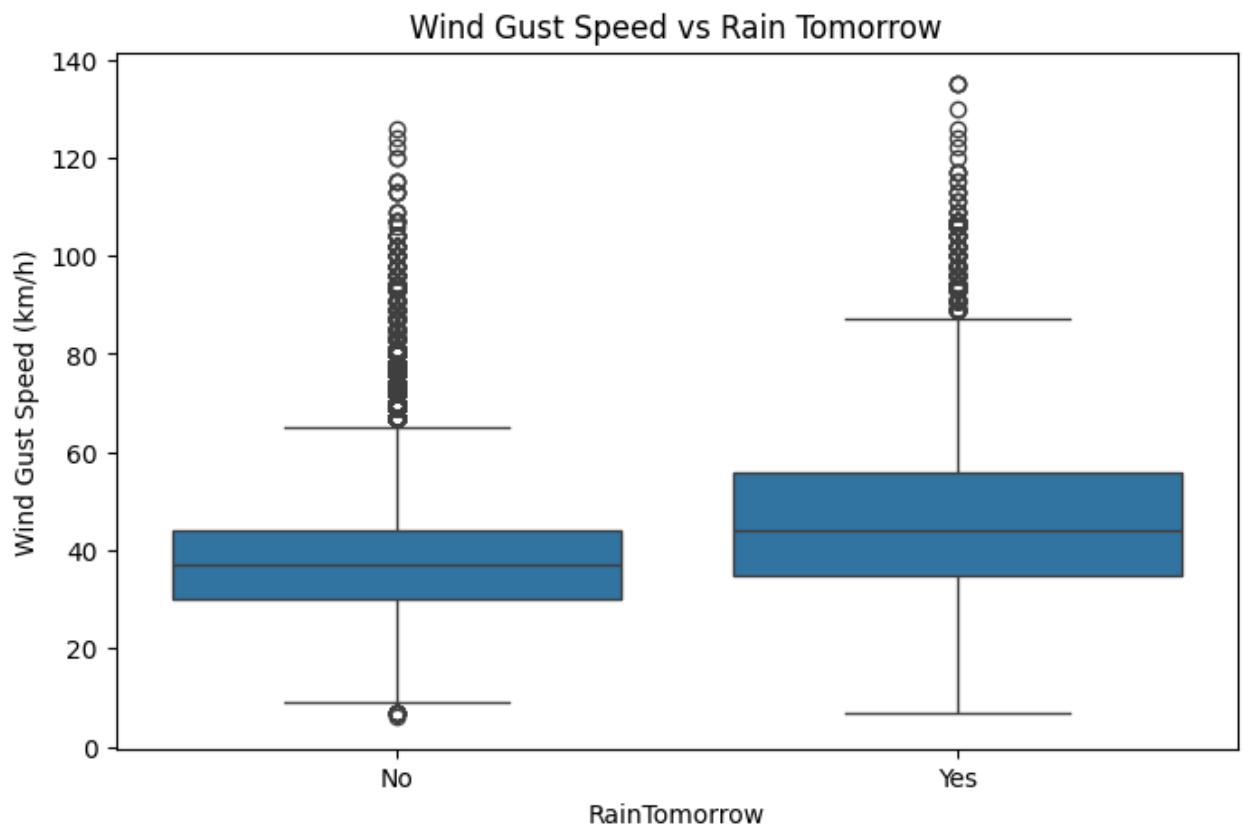
```
In [13]: plt.figure(figsize=(8,5))
sns.boxplot(
    data=rainAUS,
    x="RainTomorrow",
    y="Pressure3pm"
)
plt.title("Afternoon Pressure vs Rain Tomorrow")
plt.ylabel("Pressure at 3pm (hPa)")
plt.show()

# Rainy days tomorrow are preceded by lower pressure. This matches meteor
# Falling pressure is an early sign of incoming rain.
```

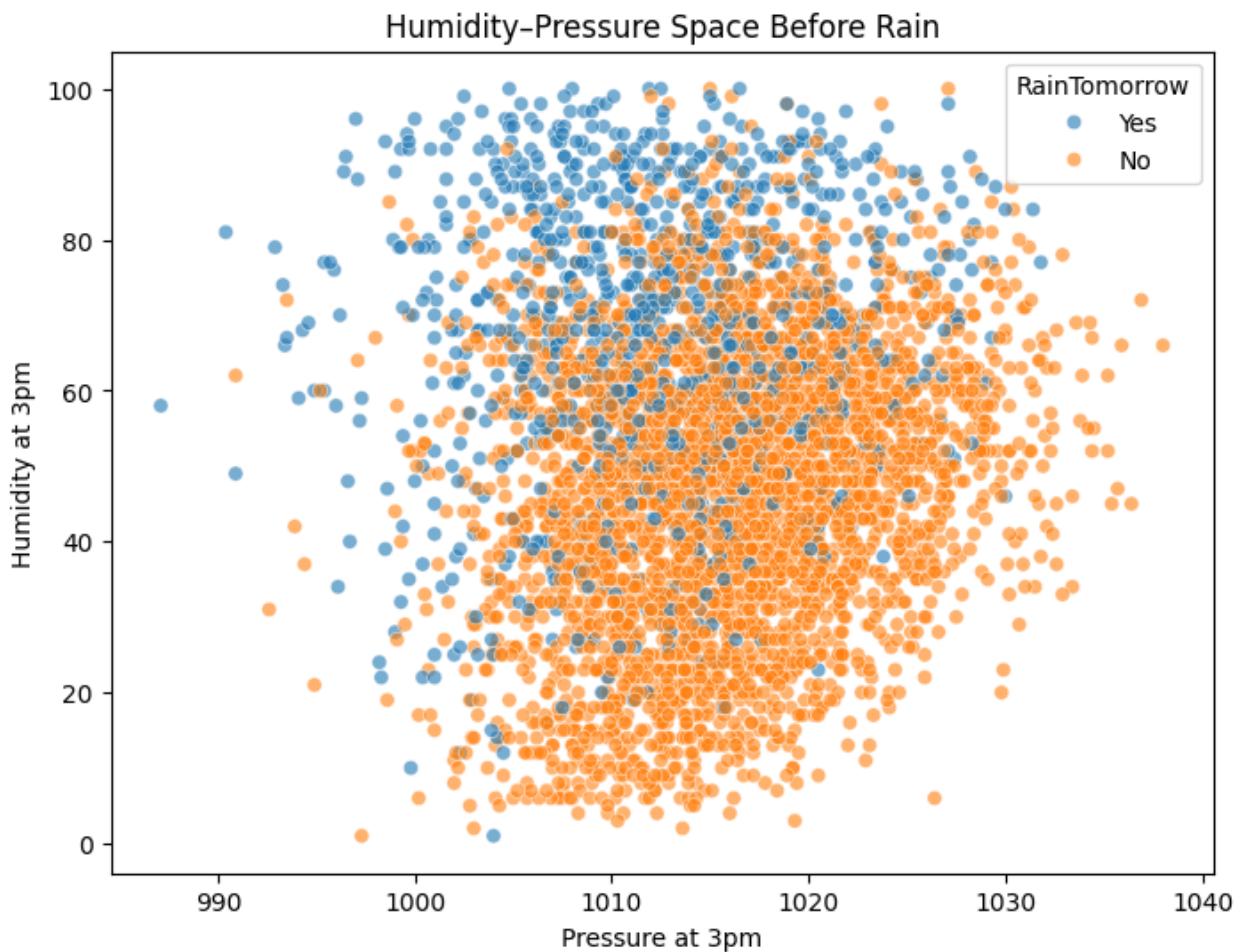


```
In [14]: plt.figure(figsize=(8,5))
sns.boxplot(
    data=rainAUS,
    x="RainTomorrow",
    y="WindGustSpeed"
)
plt.title("Wind Gust Speed vs Rain Tomorrow")
plt.ylabel("Wind Gust Speed (km/h)")
plt.show()

# Wind strengthens before rain. Extreme values are more frequent.
```



```
In [15]: plt.figure(figsize=(8,6))
sns.scatterplot(
    data=rainAUS.sample(5000, random_state=42),
    x="Pressure3pm",
    y="Humidity3pm",
    hue="RainTomorrow",
    alpha=0.6
)
plt.title("Humidity–Pressure Space Before Rain")
plt.xlabel("Pressure at 3pm")
plt.ylabel("Humidity at 3pm")
plt.show()
```



```
In [16]: # Dataset is divided into input features ('X') and the target variable ('y')
# All explanatory variables are assigned to X, while 'RainTomorrow' is isolated as y

X = rainAUS.drop(columns="RainTomorrow")
y = rainAUS["RainTomorrow"]
```

```
In [17]: # Target variable encoding - 'RainTomorrow' is converted into binary numerical values
# The labels *No* and *Yes* are mapped to 0 and 1 - where 1 indicates the occurrence of rain

y = y.map({"No": 0, "Yes": 1})
```

```
In [18]: # New features representing the year, month, and day are derived from the Date column
# potential seasonal and time-related patterns in rainfall occurrence. These new features are added to X

X["Date"] = pd.to_datetime(X["Date"])
X["Year"] = X["Date"].dt.year
X["Month"] = X["Date"].dt.month
X["Day"] = X["Date"].dt.day

X = X.drop(columns="Date")
```

```
In [19]: # Separating numerical and categorical variables

from sklearn.impute import SimpleImputer
num_cols = X.select_dtypes(include="number").columns
```

```
cat_cols = X.select_dtypes(include="object").columns
```

```
In [20]: # Handling missing values - categorical variables are imputed using the mode
# No missing values - suitable for model training.
```

```
cat_imputer = SimpleImputer(strategy="most_frequent")
X[cat_cols] = cat_imputer.fit_transform(X[cat_cols])
```

```
num_imputer = SimpleImputer(strategy="median")
X[num_cols] = num_imputer.fit_transform(X[num_cols])
```

```
In [21]: #One-hot encoding
```

```
# categorical features are transformed into numerical format using one-hot encoding
# allowing machine learning algorithms to process categorical information
```

```
X = pd.get_dummies(X, columns=cat_cols, drop_first=True)
print(X.head())
```

```

      MinTemp  MaxTemp  Rainfall  WindGustSpeed  WindSpeed9am  WindSpeed3pm
\ 
0    13.4     22.9     0.6        44.0       20.0       24.0
1     7.4     25.1     0.0        44.0        4.0       22.0
2    12.9     25.7     0.0        46.0       19.0       26.0
3     9.2     28.0     0.0        24.0       11.0        9.0
4    17.5     32.3     1.0        41.0        7.0      20.0

  Humidity9am  Humidity3pm  Pressure9am  Pressure3pm  ...  WindDir3pm_NW
\ 
0     71.0      22.0     1007.7     1007.1   ...
1     44.0      25.0     1010.6     1007.8   ...
2     38.0      30.0     1007.6     1008.7   ...
3     45.0      16.0     1017.6     1012.8   ...
4     82.0      33.0     1010.8     1006.0   ...
                                         ...      True

  WindDir3pm_S  WindDir3pm_SE  WindDir3pm_SSE  WindDir3pm_SSW  WindDir3pm
_SW \
0     False      False      False      False      False      Fa
lse
1     False      False      False      False      False      Fa
lse
2     False      False      False      False      False      Fa
lse
3     False      False      False      False      False      Fa
lse
4     False      False      False      False      False      Fa
lse

  WindDir3pm_W  WindDir3pm_WNW  WindDir3pm_WSW  RainToday_Yes
0     False      True      False      False
1     False      False      True      False
2     False      False      True      False
3     False      False      False      False
4     False      False      False      False

```

[5 rows x 109 columns]

```
In [22]: # All variables are now bool or float64 type. There are no missing values

print("Frequency of data type")
print(X.dtypes.value_counts())

print(y.isna().sum())
```

```
Frequency of data type
bool      94
float64   15
Name: count, dtype: int64
0
```

```
In [23]: # TRAIN-TEST SPLIT and DECISION TREE MODEL TRAINING
# Data is split into training and test set (80% of the observations used)
```

```
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, shuffle=True
)

# A Decision Tree classifier is then initialized with constraints on tree

tree = DecisionTreeClassifier(
    max_depth=5,
    min_samples_leaf=50,
    random_state=42
)

tree.fit(X_train, y_train)
```

Out[23]:

```
▼ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5, min_samples_leaf=50, random_state=42)
```

In [24]:

```
# Model evaluation and generating predictions on the test set. Model perf
#
```

```
from sklearn.metrics import classification_report, confusion_matrix

y_pred = tree.predict(X_test)

print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
[[21090  928]
 [ 3730 2691]]
```

	precision	recall	f1-score	support
0	0.85	0.96	0.90	22018
1	0.74	0.42	0.54	6421
accuracy			0.84	28439
macro avg	0.80	0.69	0.72	28439
weighted avg	0.83	0.84	0.82	28439

In [25]:

```
# Re-trained with class weights set to balances in order to address the imb
# Re-evaluating on the test set using same metrics.
```

```
tree = DecisionTreeClassifier(
    max_depth=5,
    min_samples_leaf=50,
```

```
    class_weight="balanced",
    random_state=42
)
tree.fit(X_train, y_train)

from sklearn.metrics import classification_report, confusion_matrix

y_pred = tree.predict(X_test)

print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))

# This model misses far fewer rainy days, which is the main objective of
# The balanced Decision Tree provides a more appropriate model for rainfa
```

```
[[17266 4752]
 [ 1774 4647]]
```

	precision	recall	f1-score	support
0	0.91	0.78	0.84	22018
1	0.49	0.72	0.59	6421
accuracy			0.77	28439
macro avg	0.70	0.75	0.71	28439
weighted avg	0.81	0.77	0.78	28439

```
In [26]: # RANDOM FOREST MODEL TRAINING AND EVALUATION
```

```
# Class weights are set to balanced to address the class imbalance and to
# prioritise correct identification of rainy days. Hyperparameters such as
# and minimum leaf size are chosen to control model complexity and reduce
# The trained model is evaluated on the test set using a confusion matrix
```

```
from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(
    n_estimators=300,
    max_depth=10,
    min_samples_leaf=50,
    class_weight="balanced",
    random_state=42,
    n_jobs=-1
)

rf.fit(X_train, y_train)

from sklearn.metrics import classification_report, confusion_matrix

y_pred_rf = rf.predict(X_test)

print(confusion_matrix(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))
```

	precision	recall	f1-score	support
0	0.91	0.80	0.85	22018
1	0.51	0.74	0.61	6421
accuracy			0.78	28439
macro avg	0.71	0.77	0.73	28439
weighted avg	0.82	0.78	0.80	28439

```
In [27]: # GRADIENT BOOSTING MODEL TRAINING AND EVALUATION
# Model performance is first evaluated using the default classification t
# probabilities are used to adjust the decision threshold to 0.3 in order
# improve recall for the rainfall class. The impact of this threshold adj

from sklearn.ensemble import GradientBoostingClassifier

gb = GradientBoostingClassifier(
    n_estimators=200,
    learning_rate=0.05,
    max_depth=3,
    random_state=42
)

gb.fit(X_train, y_train)

from sklearn.utils.class_weight import compute_sample_weight

sample_weights = compute_sample_weight(
    class_weight="balanced",
    y=y_train
)

gb.fit(X_train, y_train, sample_weight=sample_weights)

y_pred_gb = gb.predict(X_test)

print(confusion_matrix(y_test, y_pred_gb))
print(classification_report(y_test, y_pred_gb))

y_prob_gb = gb.predict_proba(X_test)[:, 1]
y_pred_gb = (y_prob_gb > 0.3).astype(int)

print(classification_report(y_test, y_pred_gb))
```

```
[ [17674  4344]
 [ 1619  4802]]
      precision    recall   f1-score   support
      0         0.92     0.80     0.86    22018
      1         0.53     0.75     0.62     6421

      accuracy          0.79    28439
      macro avg       0.72     0.78     0.74    28439
  weighted avg       0.83     0.79     0.80    28439

      precision    recall   f1-score   support
      0         0.95     0.59     0.73    22018
      1         0.39     0.90     0.54     6421

      accuracy          0.66    28439
      macro avg       0.67     0.74     0.64    28439
  weighted avg       0.82     0.66     0.69    28439
```

In [28]: # COMPARISON

```
models = {
    "Decision Tree": tree,
    "Random Forest": rf,
    "Gradient Boosting": gb
}

for name, model in models.items():
    y_pred = model.predict(X_test)
    print(
        name,
        "Recall (rain):", recall_score(y_test, y_pred),
        "F1 (rain):", f1_score(y_test, y_pred)
    )
```

```
Decision Tree Recall (rain): 0.7237190468774334 F1 (rain): 0.5874841972187
105
Random Forest Recall (rain): 0.7375798162280018 F1 (rain): 0.6064797029069
023
Gradient Boosting Recall (rain): 0.7478585890048279 F1 (rain): 0.616946103
9378171
```

In [29]:

```
models = {
    "Decision Tree": tree,
    "Random Forest": rf,
    "Gradient Boosting": gb
}

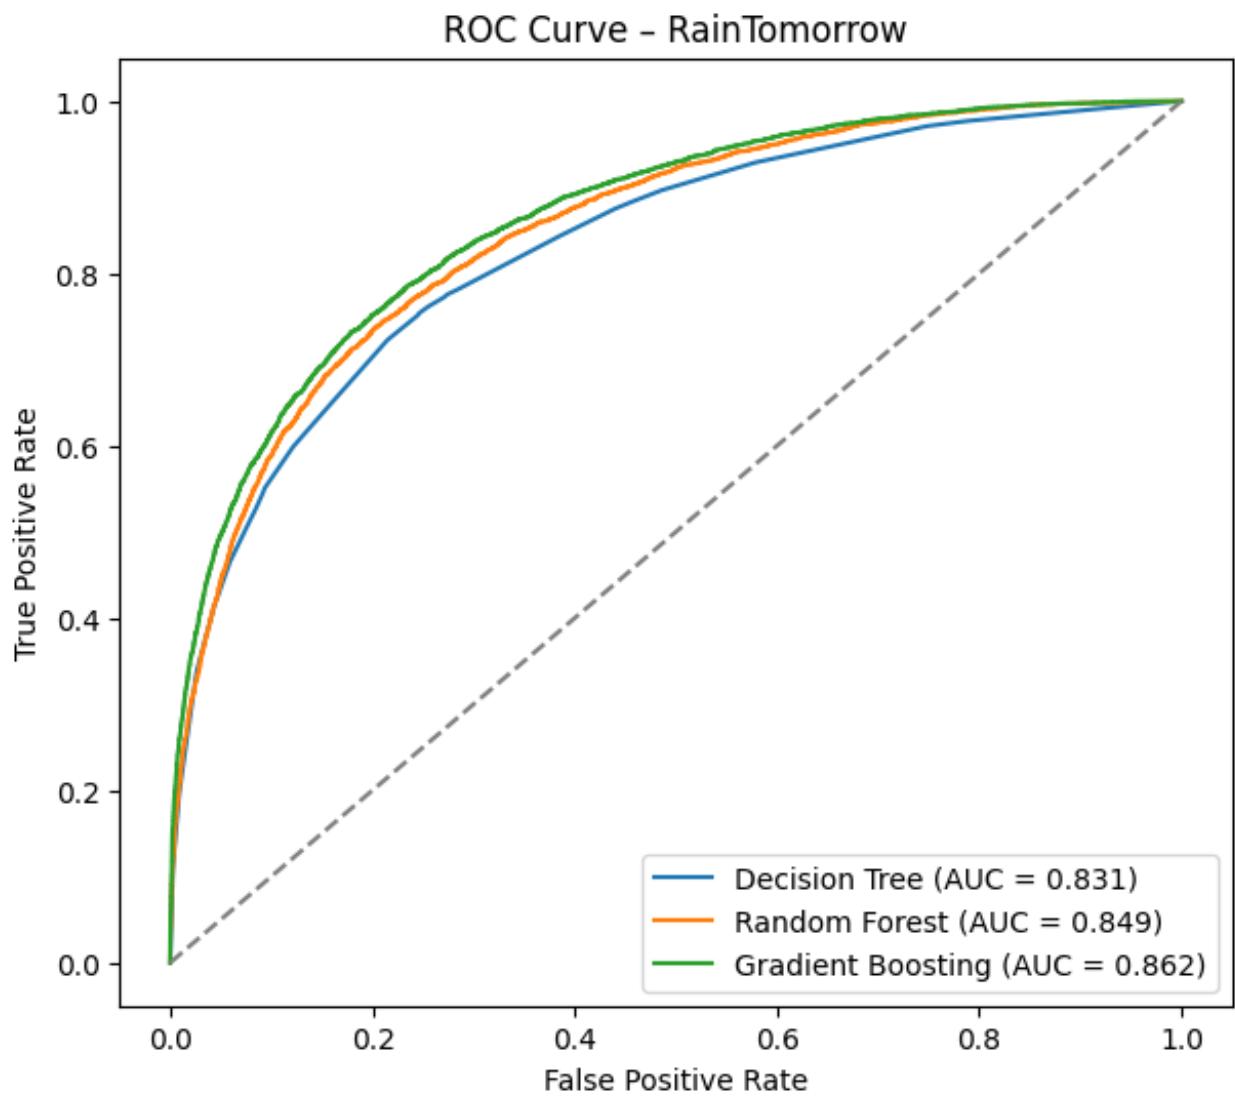
plt.figure(figsize=(7,6))

for name, model in models.items():
```

```
y_prob = model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.3f})')

plt.plot([0,1], [0,1], linestyle="--", color="gray")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve – RainTomorrow")
plt.legend()
plt.show()

# the ROC analysis confirms that ensemble methods outperform a single tree
# reliable model for rainfall prediction, particularly when probability-based
```



```
In [30]: # Comparing the classification performance of a Decision Tree, Random Forest, and Gradient Boosting

# Decision Tree
y_pred_dt = tree.predict(X_test)

# Random Forest
y_pred_rf = rf.predict(X_test)
```

```
# Gradient Boosting (threshold = 0.3)
y_prob_gb = gb.predict_proba(X_test)[:, 1]
y_pred_gb = (y_prob_gb > 0.3).astype(int)
```

```
In [31]: def plot_confusion(ax, y_true, y_pred, title):
    cm = confusion_matrix(y_true, y_pred)
    cm_pct = cm / cm.sum() * 100

    labels = np.array([
        f"TN\n{cm[0,0]}:{},\n{cm_pct[0,0]:.1f}%",
        f"FP\n{cm[0,1]}:{},\n{cm_pct[0,1]:.1f}%",
        f"FN\n{cm[1,0]}:{},\n{cm_pct[1,0]:.1f}%",
        f"TP\n{cm[1,1]}:{},\n{cm_pct[1,1]:.1f}%"
    ])

    sns.heatmap(
        cm,
        annot=labels,
        fmt="",
        cmap="Greens",
        cbar=False,
        ax=ax
    )
    ax.set_title(title)
    ax.set_xlabel("Predicted")
    ax.set_ylabel("Actual")
```

```
In [32]: fig, axes = plt.subplots(1, 3, figsize=(18, 5))

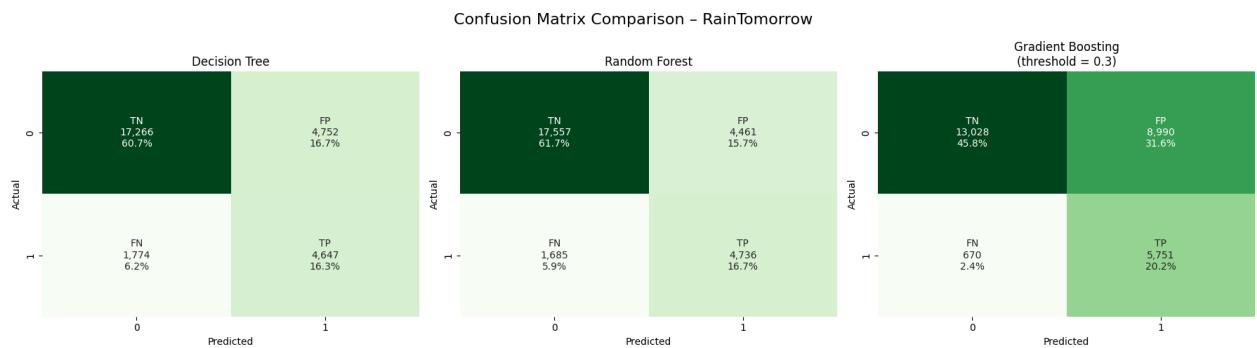
plot_confusion(
    axes[0],
    y_test,
    y_pred_dt,
    "Decision Tree"
)

plot_confusion(
    axes[1],
    y_test,
    y_pred_rf,
    "Random Forest"
)

plot_confusion(
    axes[2],
    y_test,
    y_pred_gb,
    "Gradient Boosting\n(threshold = 0.3)"
)

plt.suptitle("Confusion Matrix Comparison – RainTomorrow", fontsize=16)
plt.tight_layout()
```

```
plt.show()
```

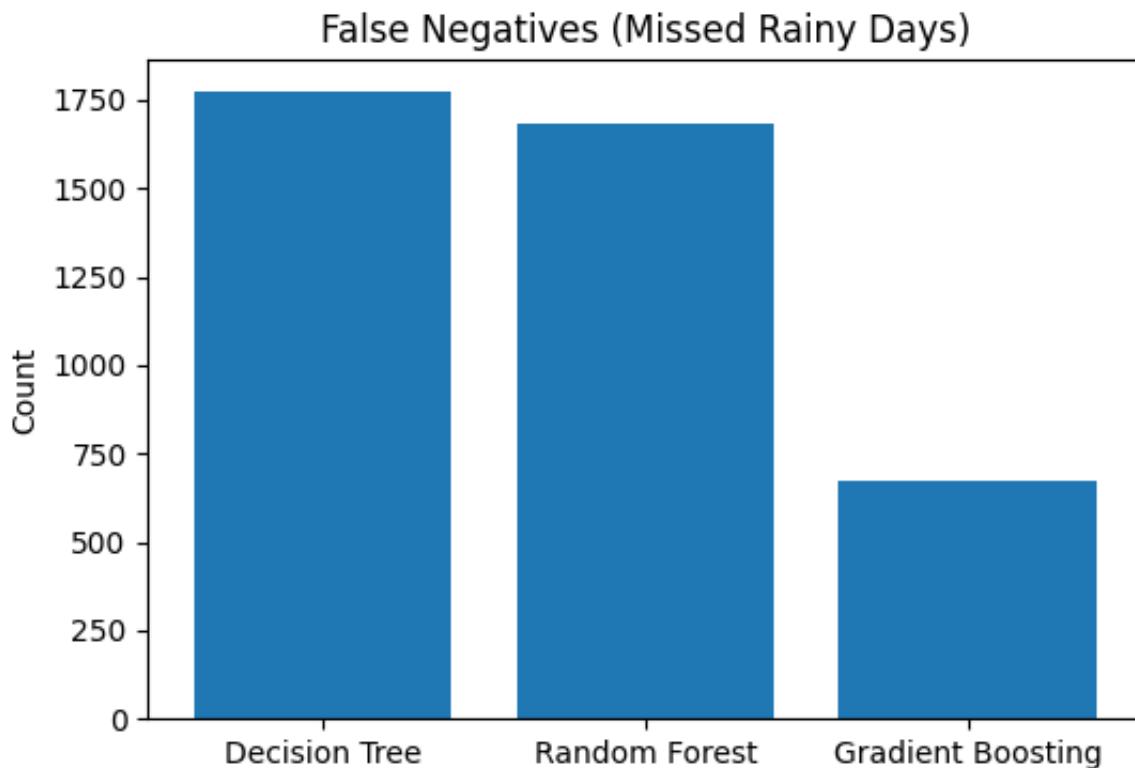


```
In [33]: # FN comparison bar chart
```

```
fn_counts = {
    "Decision Tree": confusion_matrix(y_test, y_pred_dt)[1,0],
    "Random Forest": confusion_matrix(y_test, y_pred_rf)[1,0],
    "Gradient Boosting": confusion_matrix(y_test, y_pred_gb)[1,0]
}

plt.figure(figsize=(6,4))
plt.bar(fn_counts.keys(), fn_counts.values())
plt.title("False Negatives (Missed Rainy Days)")
plt.ylabel("Count")
plt.show()

# visually says that Gradient Boosting is the best
# why?
# False Negatives = model is predicting "No Rain", but in reality it rain
# Missing rain is far more costly than prediciting the rain that does not
# In rainfall prediction, false negatives are the most critical error bec
# to predict rain leads to higher real-world costs and risks than issuing
```



```
In [ ]: # CONCLUSIONS
# The ROC curves show that ensemble methods (Random Forest and Gradient B
# False Negative analysis highlights that Gradient Boosting produces the
# The Gradient Boosting model achieves the best trade-off between precisi
#
# The Gradient Boosting Classifier is the most suitable model for rainfall
#It provides:
#
#The best overall predictive performance (highest ROC-AUC and F1-score).
#The best recall for rain predictions (fewest missed rainy days).
#Robustness against class imbalance when threshold tuning and sample weig
#Thus, it is the recommended model for deployment in a rainfall predictio
```